

fixed lr (LoRA 3e-4, embed 1e-3): bar = student hit rate (mean over seeds), dots = seeds, gray = base model (no fine-tuning)

base model (no fine-tuning) LoRA r8, stock prompt embedding matrix only, stock prompt

