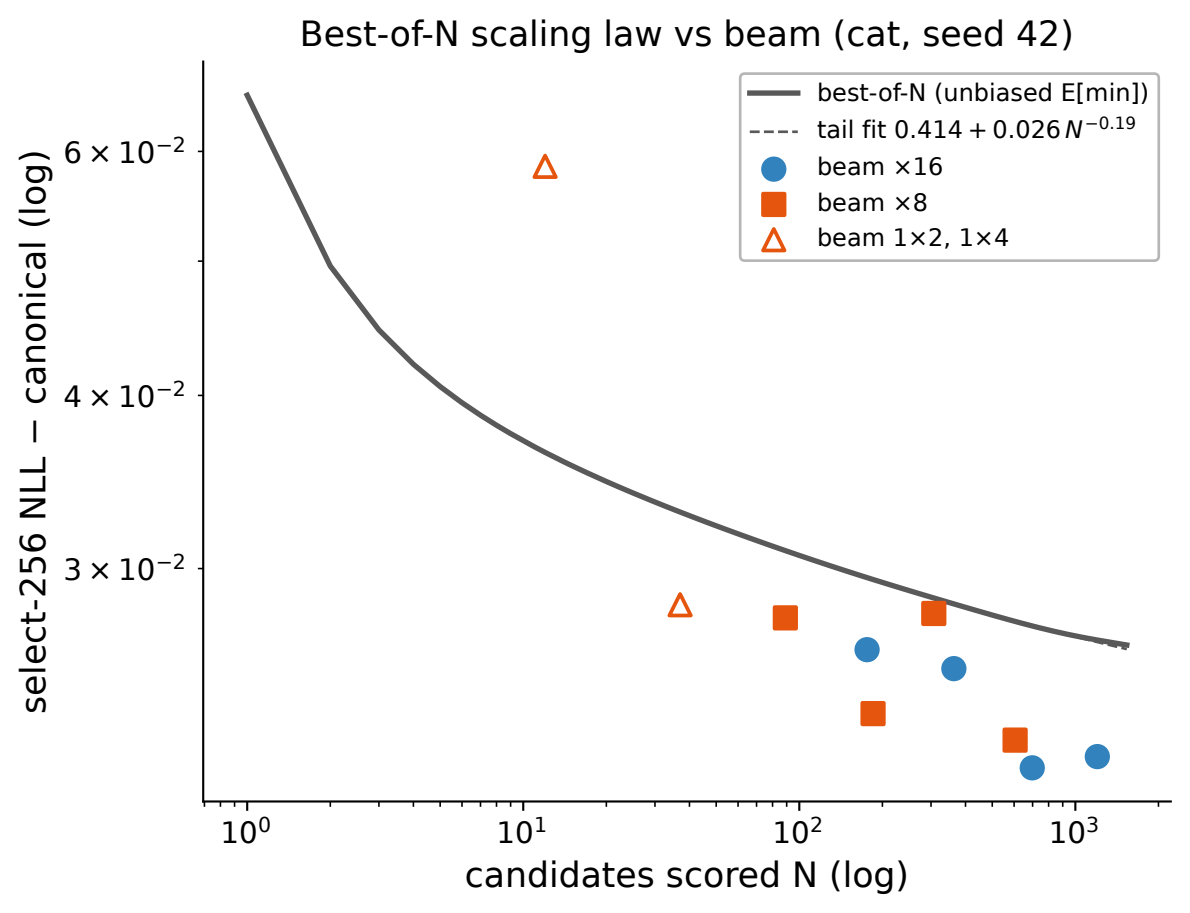
Best-of-N scaling law vs beam (cat, seed 42)
best-of-N (unbiased E[min])
tail fit $0.414 + 0.026\,N^{-0.19}$
beam ×16
beam ×8
beam 1×2, 1×4
select-256 NLL − canonical (log)
6×10^{-2}
4×10^{-2}
3×10^{-2}
10^{0}
10^{1}
10^{2}
10^{3}
candidates scored N (log)