

Same 2500-step latent budget: interleave K verbalizations, or decode at the end? (cat, seed 42)

